

Flight Delay Big Data Analysis - Draft Final Report

Luca Foresti

2026-05-20

Contents

1	Executive Summary	1
2	Dataset And Preparation	2
2.1	Canonical Fields	2
2.2	Cleaning Policy	2
2.3	Generated Input Sizes	3
3	Analyses And Implementations	3
3.1	Analysis 1: Delay By Airport, Month, And Delay Range	3
3.2	Analysis 2: Airline-Airport Ranking	3
3.3	Spark SQL	4
3.4	Spark Core	4
3.5	Hive	4
4	Produced Result Samples	4
4.1	Delay Sample, First 10 Rows	4
4.2	Ranking Sample, First 10 Rows	5
5	Benchmark Evidence	5
5.1	Benchmark Pivot	5
5.2	Benchmark Charts	6
6	Critical Comparison	7
7	Scalability And Execution Setting	8
8	Reproducibility And Validation	8
9	Limitations	8
10	Conclusion	9

1 Executive Summary

This project analyzes the 2024 United States flight delay dataset with a reproducible big-data workflow. The repository prepares a canonical cleaned Parquet dataset, implements two analytical jobs with three technologies, records execution-time evidence, and produces report-ready tables and charts.

Repository: <https://github.com/Forest904/flight-delay-big-data-analysis.git>

Spark SQL is used as the correctness reference because its declarative grouping and window functions make the required analyses easiest to express and review. Spark Core reimplements the same logic with RDD trans-

formations. Hive is the third required technology and runs through a Docker-based Hive stack. MapReduce is intentionally out of scope for this submission.

The two implemented analyses are:

- delay report by departure airport, month, and delay range;
- airline-airport ranking by average departure delay.

The main reproducibility commands are:

```
make setup
make check-env
make prepare
make generate-sizes
make run-spark-sql
make run-spark-core
make run-hive
make benchmark-local
make benchmark-cluster
make charts
make report
```

2 Dataset And Preparation

The source dataset is Kaggle [hrishitpatil/flight-data-2024](#), stored locally as `data/raw/flight_data_2024.csv`. The raw CSV has 7,079,081 rows, 35 columns, and a local file size of 1,309,010,752 bytes. The raw file is not committed to Git because it is large and externally hosted.

The preparation pipeline is implemented by `src/preparation/prepare_spark.py` and is run with:

```
make prepare
```

The prepared output is `data/prepared/flights_2024_clean.parquet`. It is the canonical input for Spark SQL, Spark Core, and Hive.

2.1 Canonical Fields

The prepared dataset keeps the fields required by the selected analyses:

- Date and grouping: `flight_date`, `month`.
- Airline fields: `airline_code`, plus nullable `airline_name` because the raw dataset has no airline-name field.
- Airport fields: `origin_airport`, `destination_airport`.
- Delay metrics: `departure_delay`, `arrival_delay`.
- Cancellation and diversion: `cancelled`, `diverted`, `cancellation_code`.
- Delay causes: carrier, weather, NAS, security, and late-aircraft delay minutes.

2.2 Cleaning Policy

The cleaning policy is conservative. Rows are removed only when structural fields are invalid after casting: missing or invalid flight date, month, airline code, origin airport, destination airport, cancelled flag, or diverted flag. Cancelled flights, diverted flights, null delay values, and negative delay values are preserved.

Empty strings are normalized to null. Numeric delay fields are cast to numeric types but are not imputed. This matters because Spark and Hive averages ignore null values, while cancellation-rate calculations still need cancelled flights in the denominator.

The raw data inspection found 92,970 null departure-delay rows and 113,814 null arrival-delay rows. These rows are retained unless they fail a structural rule. Negative delay values are preserved because they represent early departures or arrivals.

2.3 Generated Input Sizes

The input-size generator creates controlled benchmark inputs under `data/generated/`:

- 100k
- 500k
- 1m
- 3m
- full

Smaller inputs are selected with a deterministic hash-based method using seed 20240520, avoiding chronological bias from simply taking the first rows. Optional larger inputs, such as 14m and 28m, can be generated by controlled replication when explicitly requested, but they were not required for the current report artifacts.

3 Analyses And Implementations

3.1 Analysis 1: Delay By Airport, Month, And Delay Range

This job groups flights by departure airport, month, and delay range:

- low: `departure_delay < 15`
- medium: `15 <= departure_delay <= 60`
- high: `departure_delay > 60`

Rows with null departure delay are excluded from this job because they cannot be assigned to one of the required delay ranges. For each group, the job computes flight count, average departure delay, average arrival delay, and the most frequent delay or cancellation cause.

Textual pseudocode:

```
filter rows where departure_delay is not null
derive delay_range from departure_delay
derive cause from cancellation code or largest positive delay-cause field
group by origin_airport, month, delay_range
compute counts and averages
count causes inside each group
choose top cause by count desc, cause label asc
order output deterministically
```

3.2 Analysis 2: Airline-Airport Ranking

This job ranks airlines at each departure airport by average departure delay. For each airport-airline pair it computes flight count, average departure delay, average arrival delay, cancellation rate, airport average departure delay, difference from the airport average, and rank at airport.

Textual pseudocode:

```
group flights by origin_airport and airline_code
compute airline-level counts, averages, and cancellation rate
group flights by origin_airport
compute airport-level average departure delay
join airline metrics with airport metrics
rank airlines per airport by avg_departure_delay asc, nulls last
order output by airport, rank, average delay, and airline code
```

3.3 Spark SQL

Spark SQL is the reference implementation. It expresses both analyses with DataFrame temporary views, SQL aggregations, and window functions. The airline-airport ranking uses `RANK() OVER (PARTITION BY origin_airport ORDER BY avg_departure_delay ASC NULLS LAST)`. The top-cause selection uses grouped cause counts and `ROW_NUMBER()` for deterministic tie-breaking.

Spark SQL is concise and easy to validate because grouping keys, derived fields, and ranking rules are visible in the query structure. Its main cost is that window functions and grouped aggregations still require shuffles.

3.4 Spark Core

Spark Core reimplements the Spark SQL logic with RDD transformations. It uses `reduceByKey` for the delay aggregates, cause counts, airline aggregates, and airport aggregates. It avoids `groupByKey` for aggregation because reducing partial accumulator state before the shuffle is safer and more efficient.

The ranking job joins compact aggregate RDDs by airport, then sorts per-airport airline rows to reproduce Spark SQL `RANK()` semantics. Spark Core offers more control over transformation steps, but the implementation is more verbose and easier to get wrong than Spark SQL.

3.5 Hive

Hive is the third required technology. The runner starts the Docker Compose Hive stack, creates an external table over the prepared Parquet data, executes HiveQL versions of both analyses, and exports CSV outputs with the same schemas as Spark SQL and Spark Core.

Hive is useful as a SQL-on-Hadoop comparison point, but in this project it is run as a containerized local Hive service. The Docker cluster benchmark includes Hive rows for controlled execution-setting evidence, but Hive is not a distributed Hive-on-YARN cluster in this setup.

4 Produced Result Samples

The full result files are generated under `outputs/` and ignored by Git. The first-10 rows are curated under `report/tables/`. Spark Core and Hive were validated against the Spark SQL reference, so the matching samples below are shown once per analysis to keep the PDF readable. The committed evidence files exist for Spark SQL, Spark Core, and Hive under `report/tables/`, with one delay sample and one ranking sample per technology.

4.1 Delay Sample, First 10 Rows

Origin	Mo.	Range	Flights	Avg dep.	Avg arr.	Top cause
ABE	1	high	8	231.5	221.375	delay:carrier
ABE	1	low	35	-6.286	-19.629	unknown
ABE	1	medium	4	32.75	31.25	delay:late_aircraft
ABE	2	high	3	324	335.667	delay:late_aircraft
ABE	2	low	40	-8.5	-23.575	unknown
ABE	2	medium	4	45.5	29.333	unknown
ABE	3	high	8	127	121.625	delay:carrier
ABE	3	low	41	-4.756	-16.39	unknown
ABE	3	medium	3	32.333	12.667	unknown
ABE	4	high	6	158.333	150.333	delay:carrier

4.2 Ranking Sample, First 10 Rows

This output is split into two narrow tables. Together they represent the same 10 rows from the produced ranking result.

Origin	Air.	Flights	Avg dep.	Avg arr.	Cancel rate
ABE	9E	140	10.234	0.445	0.021
ABE	G4	260	12.463	4.337	0.012
ABE	OH	159	19.38	6.867	0.006
ABE	OO	40	25.297	14.389	0.075
ABI	MQ	249	6.664	4.405	0.008
ABQ	DL	218	1.383	-5.15	0.018
ABQ	OO	547	3.617	-0.855	0.004
ABQ	UA	273	4.584	-1.477	0.022
ABQ	MQ	127	7.134	5.575	0
ABQ	AS	67	7.522	1.881	0

Origin	Air.	Airport avg dep.	Diff. vs airport	Rank
ABE	9E	14.606	-4.373	1
ABE	G4	14.606	-2.143	2
ABE	OH	14.606	4.774	3
ABE	OO	14.606	10.691	4
ABI	MQ	6.664	0	1
ABQ	DL	8.792	-7.409	1
ABQ	OO	8.792	-5.176	2
ABQ	UA	8.792	-4.208	3
ABQ	MQ	8.792	-1.659	4
ABQ	AS	8.792	-1.27	5



5 Benchmark Evidence



The benchmark runner records technology, job name, input size, environment, cluster-size label, duration, output rows, status, and timestamp. The report-ready summary is generated under `report/tables/`.

5.1 Benchmark Pivot

Env.	Input	Records	Job	SQL s	Core s	Hive s
docker	100k	100000	ranking	8.676	3.693	14.976
docker	100k	100000	delay	15.454	4.164	17.509
docker	1m	1000000	ranking	7.54	3.702	14.15
docker	1m	1000000	delay	16.15	4.082	19.66
local	100k	100000	ranking	1.842	0.765	14.831
local	100k	100000	delay	8.581	0.959	18.071

5.2 Benchmark Charts





6 Critical Comparison

Spark SQL is the most maintainable implementation. Its SQL syntax makes the grouping, cause selection, and airport-partitioned ranking explicit. The tradeoff is that the engine plans shuffles for grouped aggregates and window functions, so concise code does not mean zero distributed cost.

Spark Core is the most explicit implementation. It exposes each aggregation, join, sort, and ranking step. In the available benchmark rows it is the fastest technology, especially for the local 100k runs. That result should be read with care: the RDD implementation writes small final result tables locally after aggregation, and the measured workload is not a full production cluster run. Still, Spark Core shows how map-side combining with `reduceByKey` can reduce shuffle volume compared with moving raw records by key.

Hive is the slowest technology in these runs, but it remains valuable as a third required big-data stack and a SQL-on-Hadoop comparison. The Hive runner pays service and query overhead that Spark does not pay in the same way. Its implementation is clear for SQL-style aggregation, but the local Docker setup does not demonstrate the strengths of a real distributed Hive deployment.

The delay-by-airport-month job tends to be more expensive than the airline-airport ranking in the current evidence because it derives delay ranges, counts causes, and performs top-cause selection per airport-month-range. The ranking job still requires grouping and per-airport ranking, but it operates on compact airport-airline aggregates after the first stage.

7 Scalability And Execution Setting

The project includes local benchmark evidence and Docker execution-setting evidence. The Spark Docker cluster uses a standalone Spark master and two Spark workers. This varies the execution environment, but all containers run on one physical laptop, share Docker Desktop CPU, memory, and disk limits, and read from a bind-mounted local directory rather than HDFS or object storage.

Hive benchmark rows are included in the Docker cluster benchmark output so the tables cover all three technologies. However, Hive remains a single-node containerized Hive setup in this project. It should not be presented as a distributed Hive cluster.

Because the benchmark artifacts currently include local 100k rows and Docker 100k/1m rows, the report can support cautious scalability discussion but not strong production-scaling claims. The generated input-size machinery supports larger local experiments where hardware allows.

8 Reproducibility And Validation

The repository keeps generated large data and raw data out of Git. The committed evidence is limited to small report artifacts: first-10 tables, benchmark tables, and chart images.

Correctness validation is available through:

```
.\.venv\Scripts\python.exe scripts\validate_spark_sql_outputs.py
.\.venv\Scripts\python.exe scripts\validate_spark_core_outputs.py
.\.venv\Scripts\python.exe scripts\validate_hive_outputs.py
```

Spark Core and Hive validators compare their outputs against the Spark SQL reference, including output columns, row counts, key sets, numeric values within tolerance, top-cause values, ranking order, and first-10 sample files.

9 Limitations

- The raw dataset and generated benchmark inputs are not committed to Git because they are large.
- The raw dataset does not include airline names, so analyses use `airline_code`.
- MapReduce is out of scope.
- The Docker Spark cluster is a simulation on one physical machine, not a true multi-machine cluster.
- Hive is containerized locally and is not running on a distributed Hadoop/YARN cluster.
- Windows Spark runs require care around Java, Hadoop `winutils.exe`, and native file handling; the project includes compatibility helpers and Docker paths where needed.

- Benchmark results are hardware-dependent and should be interpreted as evidence for this controlled environment, not universal performance numbers.

10 Conclusion

The project satisfies the main assignment requirements with a reproducible data preparation pipeline, two completed analyses, three technology implementations, sample output rows, benchmark tables, execution-time charts, and an honest discussion of scalability and limitations. Spark SQL provides the clearest reference logic, Spark Core demonstrates lower-level distributed transformation control, and Hive provides the required SQL-on-Hadoop comparison point.